

Katie Wood

CSE 544 Winter 2024

Implementing SimpleDB

I strove for simplicity in my implementation. In `BufferPool.java`, I used a first-in first-out (FIFO) eviction policy. While imperfect as far as database performance, a FIFO policy was the easiest for me to implement. One justification for a FIFO policy is that a user is likely to run several queries in a row on roughly the same data. For instance, when writing a complex query, I find it helpful to test several simpler queries along the way. The span of data I am accessing during this testing process is roughly constant. Once I finish, if I begin writing a new complex query, I will then access a new span of data roughly the same size as but distinct from that of complex query #1. A FIFO policy anticipates this kind of workflow since the data I have loaded most recently (say, span #1 above) stays in the buffer pool for as long as I need to hammer out complex query #1, then once I start writing complex query #2, span #1 is (eventually) evicted and replaced by span #2.

One additional point on the buffer pool: I used a `ConcurrentHashMap` for the `BufferPool` with a `ConcurrentLinkedQueue` to store eviction order. The dictionary-like structure of a hash map I found to be convenient for accessing the pages of the buffer pool given their pageids. I would not have known, however, to use a `ConcurrentHashMap` (safe for multi-threading, which is our use-case in `BufferPool`), without taking the hint from the unused import `java.util.concurrent.ConcurrentHashMap`. I wish some mention of hash maps, concurrent hash maps, and multi-threading had been made in the coding review.

Next, throughout my implementation of SimpleDB, I added checks on the arguments being passed into functions to ensure that the functions would work properly. For example, in the `isSlotUsed(int i)` method of `HeapPage`, I required that $0 \leq i < \text{numSlots}$. This prevents our erroneous use of the high-order bits of the last header byte. As another example, in the `hasNext()` method of the `HeapFileIterator` within `HeapFile`, I first checked whether the current iterator was null, and if so returned false.

This was key to ensuring that the iterator was first initialized via the `open()` call before the rest of `hasNext()` was allowed to run.

Lastly, the Join operator I found especially challenging and interesting. I opted for a nested loops join, again for the sake of simplicity. To condense my implementation of `fetchNext()`, I wrote a helper function `joinTuples(Tuple t1, Tuple t2)` for concatenating pairs of tuples that pass the join predicate. One final challenge in writing `fetchNext()` was to ensure that `next()` was not called on `child1` until all the tuples in `child2` had been scanned. To do this, I added the boolean variable **found**, which if true, causes `fetchNext()` to skip the `next()` call to `child1` until all tuples of `child2` have been scanned and `child2` rewound. Then `fetchNext()` calls `next()` on `child1` and rescans `child2`, checking for matches.